Exam DP-200: Implementing an Azure Data Solution – Skills Measured

The content of this exam was updated on December 4, 2019. Please continue scrolling to the red line section below to view the changes.

Implement data storage solutions (40-45%)

Implement non-relational data stores

- implement a solution that uses Cosmos DB, Data Lake Storage Gen2, or Blob storage
- implement data distribution and partitions
- implement a consistency model in Cosmos DB
- provision a non-relational data store
- provide access to data to meet security requirements
- implement for high availability, disaster recovery, and global distribution

Implement relational data stores

- configure elastic pools
- configure geo-replication
- provide access to data to meet security requirements
- implement for high availability, disaster recovery, and global distribution
- implement data distribution and partitions for Azure Synapse Analytics
- implement PolyBase

Manage data security

- implement data masking
- encrypt data at rest and in motion

Manage and develop data processing (25-30%)

Develop batch processing solutions

- develop batch processing solutions by using Data Factory and Azure Databricks
- ingest data by using PolyBase
- implement the integration runtime for Data Factory
- create linked services and datasets
- create pipelines and activities
- create and schedule triggers
- implement Azure Databricks clusters, notebooks, jobs, and autoscaling
- ingest data into Azure Databricks

Develop streaming solutions

- configure input and output
- select the appropriate windowing functions
- implement event processing by using Stream Analytics
- ingest and query streaming data with Azure Data Explorer

Monitor and optimize data solutions (30-35%)

Monitor data storage

- monitor relational and non-relational data sources
- implement Blob storage monitoring
- implement Data Lake Storage monitoring
- implement SQL Database monitoring
- implement Azure Synapse Analytics monitoring
- implement Cosmos DB monitoring
- implement Azure Data Explorer monitoring
- configure Azure Monitor alerts
- implement auditing by using Azure Log Analytics

Monitor data processing

- monitor Data Factory pipelines
- monitor Azure Databricks
- monitor Stream Analytics
- configure Azure Monitor alerts
- implement auditing by using Azure Log Analytics

Optimize Azure data solutions

- troubleshoot data partitioning bottlenecks
- optimize Data Lake Storage
- optimize Stream Analytics
- optimize Azure Synapse Analytics
- optimize SQL Database
- manage the data lifecycle

See below changes as of December 4, 2019

Implement data storage solutions (40-45%)

Implement non-relational data stores

- implement a solution that uses Cosmos DB, Data Lake Storage Gen2, or Blob storage
- implement data distribution and partitions
- implement a consistency model in Cosmos DB
- provision a non-relational data store
- provide access to data to meet security requirements
- implement for high availability, disaster recovery, and global distribution

Implement relational data stores

- configure elastic pools
- configure geo-replication
- provide access to data to meet security requirements
- implement for high availability, disaster recovery, and global distribution
- implement data distribution and partitions for Azure Synapse Analytics
- ~~implement data distribution and partitions for SQL Data Warehouse~~
- implement PolyBase

Manage data security

- implement data masking
- encrypt data at rest and in motion

Manage and develop data processing (25-30%)

Develop batch processing solutions

- develop batch processing solutions by using Data Factory and Azure Databricks
- ingest data by using PolyBase
- implement the integration runtime for Data Factory
- create linked services and datasets
- create pipelines and activities
- create and schedule triggers
- implement Azure Databricks clusters, notebooks, jobs, and autoscaling
- ingest data into Azure Databricks

Develop streaming solutions

- configure input and output
- select the appropriate windowing functions
- implement event processing by using stream analytics
- ingest and query streaming data with Azure Data Explorer

Monitor and optimize data solutions (30-35%)

Monitor data storage

- monitor relational and non-relational data sources
- implement Blob storage monitoring
- implement Data Lake Storage monitoring
- implement SQL Database monitoring
- implement Azure Synapse Analytics monitoring
- ~~implement SQL Data Warehouse monitoring~~
- implement Cosmos DB monitoring
- implement Azure Data Explorer monitoring
- configure Azure Monitor alerts
- implement auditing by using Azure Log Analytics

Monitor data processing

- monitor Data Factory pipelines
- monitor Azure Databricks
- monitor Stream Analytics
- configure Azure Monitor alerts
- implement auditing by using Azure Log Analytics

Optimize Azure data solutions

- troubleshoot data partitioning bottlenecks
- optimize Data Lake Storage
- optimize Stream Analytics
- optimize Azure Synapse Analytics
- ~~optimize SQL Data Warehouse~~
- optimize SQL Database
- manage the data lifecycle